

SHRINCS Parameter Search Methodology

May 20, 2026

Contents

1	Introduction	1
2	Parameters Description	2
2.1	Initial Bounds	2
2.2	SLH-DSA Parameters	2
3	Search Methodology	3
3.1	Initial Search Conditions	3
3.2	Toolset	3
3.3	Sweep Strategy	3
3.3.1	Case 1: Best in Class	4
3.3.2	Case 2: Balanced Metrics	4
3.3.3	Case 3: Adaptive Search	5
3.4	Decisional Weights	7

1 Introduction

In the case of Bitcoin, hash-based signature schemes are the most compatible among all quantum-resistant signature algorithms. Still, the smallest stateless hash-based signature proposed by NIST (SLH-DSA-128s) is 7,856 bytes [?], which is more than 100× larger than the current Schnorr signature. The question is whether we can have something smaller and more practical without significant deviation from the standard.

SHRINCS [?] proposes a hybrid signature consisting of a stateful and a stateless part, with signature size ranging from 324 bytes (stateful) to X bytes (stateless). This document attempts to answer the question of what the value of X might be while keeping other signature metrics (KeyGen, SigGen, SigVer, etc.) practical.

In this document, we present the methodology for searching parameters of SLH-DSA-compatible signatures: how changes in the hyper-tree structure and other parameters affect the signature metrics. As mentioned above, we deliberately stay within very conservative bounds and avoid any changes to the signature sub-algorithms in order to keep the construction as compatible as possible with the SLH-DSA standard.

Along the way, we will occasionally reference promising proposals such as hyper-tree pruning [?] and the +C compression technique [?], which allow for even better shrinking results. Although some of them can be considered partially compatible, we do not rely on them in the evaluation framework. We still think it is useful to outline the potential of such techniques and to let the reader tweak the parameters accordingly.

The main contributions of this report are (1) a sweep script, (2) the filtering pipeline, and (3) a recommended candidate tuple.

2 Parameters Description

2.1 Initial Bounds

As mentioned in the introduction, our objective is to find parameters that make the signature “better” than SLH-DSA. Before we can do so, we have to fix the strict boundaries that we will not cross:

1. We obviously do not want a signature larger than SLH-DSA - otherwise, why bother in the first place?
2. For Bitcoin, we require at least a 128-bit security level. None of the parameters we select may degrade security below this floor.
3. At least $q_s = 2^{40}$ supported signatures. SLH-DSA-128s defines the maximum number of signatures as $q_s = 2^{64}$, which is generous. For Bitcoin, we assume that $q_s = 2^{40}$ is a practical compromise, comfortably covering the foreseeable needs of on-chain operations and the Lightning Network.
4. Sub-algorithms, hash, and addressing functions remain unchanged.

2.2 SLH-DSA Parameters

SLH-DSA uses a hyper-tree of eXtended Merkle Signature Schemes (XMSS) to authenticate the underlying one-time and few-time signature keys. We formally parameterize each candidate configuration using the tuple (h, d, k, a, w) , where each variable controls a distinct geometric trade-off within the overall structure.

Hyper-Tree Geometry (h and d). The hyper-tree configuration determines the capacity and layer-traversal costs of the signature scheme.

- h is the total height of the multi-layered hyper-tree. This dictates the maximum number of few-time signature instances supported at the base layer, yielding a total capacity of 2^h unique leaf nodes.
- d is the number of independent tree layers comprising the hyper-tree.

To ensure a symmetric construction across all internal boundaries, our search pipeline requires that the total height h be perfectly divisible by the layer count d . Consequently, each individual XMSS tree layer has a uniform height h' :

$$h' = \frac{h}{d}.$$

Adjusting the ratio between h and d governs the balance between signature size and key-generation cost. Reducing the number of layers d shrinks the total number of intermediate public keys and authentication paths in the signature payload, saving valuable bytes. However, this exponentially increases the height h' of each constituent tree, requiring significantly more hash operations to compute the root node during key generation.

Few-Time Signature Layer (k and a). The lowest layer of the hyper-tree authenticates the Forest of Random Subsets (FORS) few-time signature scheme, which directly signs the message digest.

- k is the number of independent, parallel Merkle trees in the FORS layer.
- a is the height of each individual FORS tree, so that each constituent tree contains exactly $t = 2^a$ leaves.

Increasing k and a strengthens the scheme against multi-target subset-forgery attacks and allows the structure to support larger capacities safely. However, expanding these dimensions directly increases the byte footprint of the appended authentication paths.

One-Time Signature Base (w). Internal tree roots are signed using the Winternitz One-Time Signature Plus (WOTS+) scheme. The base parameter w defines the digit representation used to split the message digest during signing. Our framework evaluates candidate configurations across the standard baseline settings $w \in \{16, 32, 256\}$.

Selecting a larger base reduces the number of parallel hash chains ℓ , which in turn reduces the overall WOTS+ payload size. This spatial efficiency comes at the cost of local compute: larger base parameters require longer hash chains, resulting in higher execution costs during key and signature generation.

We sweep the ranges as follows:

$$(h, d, k, a, w) \in \{40, \dots, 52\} \times \{2, \dots, 26\} \times \{2, \dots, 25\} \times \{8, \dots, 21\} \times \{16, 32, 256\}.$$

The SLH-DSA-128s baseline tuple (63, 7, 14, 12, 16) sits outside this range and is included for comparison

3 Search Methodology

3.1 Initial Search Conditions

Searching for the best parameters requires defining the conditions that the chosen parameter set must satisfy. Consequently, those conditions must rely on metrics that depend on the parameters. We use the following:

- **Signature size:** the size of the signature in bytes;
- **Key generation cost:** the number of SHA-256 compressions during the key-generation phase;
- **Signature generation cost:** the number of compressions executed during signing;
- **Signature verification cost:** the number of compressions during verification;
- **Compressions per byte:** the ratio of verification cost to signature size.

We choose SLH-DSA-128s as the baseline for the search; this gives us an intuitive way to compare custom parameter sets against the standardized scheme on equal footing.

3.2 Toolset

The methodology described here uses a set of SageMath scripts available at <https://github.com/BlockstreamResearch/SPHINCS-Parameters>.

3.3 Sweep Strategy

We start by enumerating all parameter tuples (h, d, k, a, w) . Applying the initial bounds (security level, size $\leq$ SLH-DSA, at least 2^{40} supported signatures) reduces the search space from 41,328

candidates to 9,229 - a 78% cut. The remaining space is still too large to inspect by hand, so we need additional filters.

We introduce the following filters¹:

- **Bound on signature size:** $\text{target.size} \leq c_{\text{size}} \cdot \text{std.size}$, where c_{size} is a size coefficient ($0 \leq c_{\text{size}} \leq 1$).
- **Bound on key-generation cost:** $\text{target.KeyGen} \leq c_{\text{kg}} \cdot \text{std.KeyGen}$, where $c_{\text{kg}} > 0$.
- **Bound on signature-generation cost:** $\text{target.SigGen} \leq c_{\text{sg}} \cdot \text{std.SigGen}$, where $c_{\text{sg}} > 0$.
- **Bound on signature-verification cost:** $\text{target.SigVer} \leq c_{\text{sv}} \cdot \text{std.SigVer}$, where $c_{\text{sv}} > 0$.
- **Bound on the compressions-per-byte ratio:** $\text{target.C/B} \leq c_{\text{cb}} \cdot \text{std.C/B}$, where $c_{\text{cb}} > 0$.

These filters can be combined and parameterized in different ways to produce different result sets. Below, we consider several cases to illustrate how the choice of filter parameters affects the resulting candidate set.

3.3.1 Case 1: Best in Class

In this case, we find the most efficient signature for a single metric at a time, leaving the other metrics unbounded. The results are shown in Table 1.

Table 1: Absolute Global Extremes

Award Category	Tuple (h, d, k, a, w)	Size (B)	VS std, %	KeyGen (C)	VS std, ×	SigGen (C)	VS std, ×	SigVer (C)	ρ (C/B)
STANDARD Baseline	(63, 7, 14, 12, 16)	7872	+0.0%	2.93×10^5	1.00	2.28×10^6	1.00	2387	0.30
Smallest Signature	(50, 2, 6, 20, 256)	3424	-56.5%	1.55×10^{11}	5.29×10^5	3.10×10^{11}	1.36×10^5	4953	1.45
Fastest Keygen	(40, 10, 8, 20, 32)	7680	-2.4%	1.40×10^4	0.05	3.37×10^7	14.78	4673	0.61
Fastest Signing	(40, 8, 23, 8, 32)	7440	-5.5%	2.80×10^4	0.10	2.47×10^5	0.11	3888	0.52
Fastest Verification	(50, 2, 7, 17, 16)	3968	-49.6%	1.92×10^{10}	6.55×10^4	3.84×10^{10}	1.68×10^4	894	0.23
Best Density Ratio (ρ)	(50, 2, 23, 15, 16)	7840	-0.4%	1.92×10^{10}	6.55×10^4	3.84×10^{10}	1.68×10^4	1366	0.17

Table 1 confirms what one would suspect: each metric can be minimized impressively in isolation, but only at a brutal cost to the others, leaving these parameter sets unusable in practice. We therefore consider a different strategy.

3.3.2 Case 2: Balanced Metrics

In the next experiment, we look for “average” candidates whose metrics all stay below a uniform threshold $X \cdot \text{std}$ (the filter is applied to every metric, with the exception of the maximum signature size $\leq \text{std}$). By gradually increasing X , we obtain the following picture.

Naturally, the smaller the multiplier X , the fewer candidates survive the filter. A clear winner emerges at $X \approx 0.8840$; let us see how compelling it actually is (Table 2).

Table 2: Performance Comparison between the Standard Baseline and the Balanced Candidate

Metric	Standard Baseline	Optimal Candidate	Ratio
Parameters (h, d, k, a, w)	(63, 7, 14, 12, 16)	(40, 5, 24, 8, 16)	–
Signature Size (Bytes)	7,872	6,928	0.8801×
Key Generation Cost	292,862.0	146,430.0	0.5000×
Signing Cost	2,279,391.0	756,689.0	0.3320×
Verification Cost	2,387.0	1,857.0	0.7780×
Compression per Byte (ρ)	0.303227	0.268043	0.8840×

¹The coefficient for each filter, except the signature size, may be greater than 1 (meaning we accept some regression on that metric in exchange for gains elsewhere).




Figure 1: Global Minimax Expansion: Concentric Euler diagram illustrating the volume of candidates unlocked at expanding uniform multiplier thresholds.

One might call this a reasonable result. We are less impressed: a 12% saving on signature size is not particularly attractive on its own (the goal is to materially close the gap with Schnorr).

3.3.3 Case 3: Adaptive Search

The strategy is straightforward:

1. Decide which sacrifices you are willing to make on each metric.
2. Find the best candidate among those that remain.

For step (1) we apply the filters introduced in Section 3.3. This subsection is dedicated to step (2).

Each candidate is uniquely defined by its tuple of metric values, so the search for an optimal candidate can be cast as finding the point closest to the origin in a five-dimensional space (analogous to a shortest-vector problem, but mercifully confined to a small set of survivors), where each axis represents one of our metrics. That is almost the whole story.

To build intuition, suppose we care about only two metrics - signature size and signature generation cost. We require the signature size to be no larger than that of SLH-DSA, and we are willing to tolerate a signing cost of at most $2\times$ the standard. In this case all admissible candidates fit inside the green rectangle in Figure 2, which makes brute-force inspection easy: the best candidate is simply the point closest to the origin $(0, 0)$.

We cannot directly visualize a 5-dimensional space, but we can show the filtering process in a different way.

First, we scatter the unbounded candidate set (subject only to the 128-bit security condition) on a plane, with positions chosen randomly - the layout carries no meaning, and each point is




Figure 2: Visualization of Candidates with Prioritization of Size and Signing Cost Metrics.

uniquely identified by its parameter tuple. After initialization, we apply the signature-size filter, visualized as a coloured “sieve”. We then apply each remaining filter in turn until the last one.

The resulting set contains the candidates that survived every filter; this is the pool in which we search for the best one.

As an example, consider the following filter coefficients:

$$X_{\text{size}} = 0.85, \quad X_{\text{kg}} = 1.5, \quad X_{\text{sg}} = 2, \quad X_{\text{sv}} = 0.8, \quad X_{\text{cb}} = 1.$$

The filtering process is shown in Figure 3.

We start with 19,437 candidates. The signature-size filter trims the pool to 6,202; the next filter (KeyGen) cuts it by more than 7×. Each subsequent filter prunes further, leaving 60 final candidates for inspection.

With the final 60 candidates in hand, we are ready to pick a winner.

For each surviving candidate P , let X_i denote the ratio of the candidate’s value for metric i to the corresponding standard value, where $i \in \{\text{size, kg, sg, sv, cb}\}$. Working with ratios rather than raw values keeps the five axes on a common, dimensionless scale. The Euclidean distance from the origin to P is then

$$d(P) = \sqrt{X_{\text{size}}^2 + X_{\text{kg}}^2 + X_{\text{sg}}^2 + X_{\text{sv}}^2 + X_{\text{cb}}^2}, \quad (1)$$

and the optimal candidate is the one minimizing $d(P)$.

This treats every metric as equally important, which is rarely what we actually want. Signature size, for example, usually carries far more weight than key-generation cost in a Bitcoin context. The next subsection addresses this by attaching weights to each axis.




Figure 3: Full Parameter Filtering Process

3.4 Decisional Weights

We can do better by applying the *weighted* Euclidean distance. The motivation is that weights give us a principled way to encode the relative priority of each metric. Let $w_{\text{size}}, w_{\text{kg}}, w_{\text{sg}}, w_{\text{sv}}, w_{\text{cb}} \in \mathbb{R}^+$ denote the raw weight coefficients representing the relative importance of signature size, key generation, signing, verification, and compressions per byte, respectively.

To keep the calculation proportionally consistent regardless of the absolute scale of the chosen weights, we normalize them into coefficients p_i (where $i \in \{\text{size, kg, sg, sv, cb}\}$) such that they sum to 1:

$$p_i = \frac{w_i}{w_{\text{size}} + w_{\text{kg}} + w_{\text{sg}} + w_{\text{sv}} + w_{\text{cb}}}, \quad \text{ensuring} \quad \sum p_i = 1 \quad (2)$$

Using these normalized coefficients, we calculate the Weighted Euclidean Distance (the L_2 norm) from the origin $(0, 0, 0, 0, 0)$ to the candidate's point P . The distance $d(P)$ is defined as:

$$d(P) = \sqrt{p_{\text{size}}X_{\text{size}}^2 + p_{\text{kg}}X_{\text{kg}}^2 + p_{\text{sg}}X_{\text{sg}}^2 + p_{\text{sv}}X_{\text{sv}}^2 + p_{\text{cb}}X_{\text{cb}}^2} \quad (3)$$

As before, each X_i is the ratio of the candidate's metric to its standard counterpart, so the five axes remain on a common dimensionless scale.

To visualize the distances, we map each candidate's distance onto a colour scale: lower distances shade towards blue, higher distances towards red. As an example, consider the following two weight settings:

- Case 1: $w_{\text{size}} = w_{\text{kg}} = w_{\text{sg}} = w_{\text{sv}} = w_{\text{cb}} = 1$;

- Case 2: $w_{\text{size}} = 5$, $w_{\text{kg}} = 1$, $w_{\text{sg}} = 1.5$, $w_{\text{sv}} = 2$, $w_{\text{cb}} = 1$.

The result for the first case is shown below.




Figure 4: Visualization of Searching for the Case 1.

The parameters and metric values of the best candidate are shown in Table 3.

Table 3: Performance Comparison between the Standard Baseline and the Best Shortest Distance Candidate

Metric	Standard Baseline	Best Candidate	Relative Multiplier (X)
Parameters (h, d, k, a, w)	(63, 7, 14, 12, 16)	(40, 5, 19, 9, 16)	—
Signature Size (Bytes)	7,872	6,512	$0.8272\times$
Key Generation Cost	292,862.0	146,430.0	$0.5000\times$
Signing Cost	2,279,391.0	771,034.0	$0.3383\times$
Verification Cost	2,387.0	1,809.0	$0.7579\times$
Compression per Byte (ρ)	0.303227	0.277795	$0.9161\times$

For the second case, we have the following picture:




Figure 5: Visualization of Searching for the Case 2.

The resulting values are shown in Table 4.

Table 4: Performance Comparison between the Standard Baseline and the Best Shortest Distance Candidate

Metric	Standard Baseline	Best Candidate	Relative Multiplier (X)
Parameters (h, d, k, a, w)	(63, 7, 14, 12, 16)	(40, 5, 13, 12, 16)	—
Signature Size (Bytes)	7,872	6,176	0.7846×
Key Generation Cost	292,862.0	146,430.0	0.5000×
Signing Cost	2,279,391.0	945,124.0	0.4146×
Verification Cost	2,387.0	1,771.0	0.7419×
Compression per Byte (ρ)	0.303227	0.286755	0.9457×

Comparing the two cases, the resulting metrics differ noticeably. The first case yields more balanced costs across the board, while the second case clearly favours signature size and verification at the expense of a heavier signing operation - exactly the trade-off we asked the weights to enforce.